

Monitor LibreChat

A complete walkthrough of every screen in the monitoring and audit portal for a corporate LibreChat instance — adoption, usage, cost, and a shadow-mode security detector.

12 screens documented 2 interface languages (EN / PT) 0 required configuration to get started

00 · GETTING STARTED

Overview

The Monitor LibreChat is an independent portal that reads LibreChat's own database and shows, in one place, what today is either scattered or simply doesn't exist: who is using the tool, how much it costs, and whether any usage falls outside what's expected.

WHAT IT SHOWS

Adoption, token usage, cost, use cases by Agent, AI maturity, security detections, policies, authorized resources, audit trail, and the health of the Monitor itself.

WHAT IT DOES NOT DO

It never blocks anything. It doesn't intercept prompts. It is strictly observational — it reads the database after the fact.

WHERE THE DATA COMES FROM

Directly from LibreChat's MongoDB, through a dedicated read-only user. Without that connection, the portal shows clearly-marked example data.

i This manual was generated against an instance running with **example (synthetic) data** — that's why every screen shows the yellow "Showing example data" banner. The numbers change once connected to a real LibreChat instance; the screens and behavior are identical.

Signing in

Access uses a single username and password, set by whoever administers the Monitor (not the same login as LibreChat itself). There are no separate access profiles in this version — whoever signs in sees everything.

CORPORATE LIBRECHAT

Monitor

Adoption, cost, security, and audit dashboard.

USER

PASSWORD

Sign in

Restricted access. Every access to this panel is recorded in the audit trail.

The login screen. The footer notice exists because every access to this panel is recorded in the audit trail (section 12).

- 1 Open the Monitor's URL, provided by whoever administers the system.
- 2 Enter your **User** and **Password**.
- 3 Click **Sign in**. After eight hours of inactivity, the session expires and you'll need to sign in again.

- ! After several consecutive incorrect password attempts, the system refuses new attempts for a few minutes, even with the correct password. This is intentional — wait a bit and try again.

02 · GETTING STARTED

General navigation

After signing in, three elements appear on every screen:

Sidebar

Organized by purpose, not alphabetically: **Overview** (day-to-day — dashboard, adoption, costs), **Governance** (security, policies, alerts), and **Platform** (audit, settings, status). It can be collapsed via the "Collapse" button at the bottom, to gain screen space.

Language

The **EN** / **PT** buttons at the top switch the interface language instantly, with no page reload. The choice is saved in the browser — next time you open the Monitor, it remembers.


Auto-refresh

The whole portal refreshes itself every **60 seconds** — the "Updated Ns ago" indicator on the Dashboard shows the time since the last refresh. No manual reload needed.

- ✓ Most screens have a period selector (From / to, or shortcuts like 7d / 30d / 90d). The chosen period affects that screen's numbers only — changing the period on Costs doesn't change the Dashboard's period.

Executive Dashboard

The first screen after signing in. Brings the most important numbers together in one place, for a quick read of the overall state of AI usage in the organization.




Executive Dashboard with example data. The top cards, the hourly token chart, the per-role radar, and the numbers-by-role table.

The cards

CARD	WHAT IT MEANS
Users	Total enabled users and how many were active in the period. Clickable — opens the full user list.
Input / output / total tokens	Token volume consumed (what was sent to the model and what it answered).
Cost (USD)	Estimated cost for the period, always in dollars. See the Costs section for how this figure is computed.
Files generated / Prompts / Conversations	Raw usage volume for the period.
Conduct	A statistical indicator (not a fixed rule) comparing the most recent consumption against the historical mean. Clickable — opens the period's security detections in detail.

Clickable cards

Two cards open a more detailed window when clicked:



Clicking the "Users" card opens the full list, sorted by most recent access — with status (active/inactive), role, tokens consumed, and last-access date.

i "Last access" is the more recent of sign-in and actual tool usage. This exists because a LibreChat session lasts 7 days and isn't renewed on every access — using sign-in alone would leave this information stale.

Hourly token chart

Shows input, output, and a moving average (smooths out isolated spikes) over time. The bucket width (hourly, every 6h, or daily) adjusts itself based on the selected period, so the chart never turns into a wall of tiny bars.

Radar and role table


The radar chart compares usage profile across roles (tokens, prompts, conversations, active days, distinct models and Agents — all normalized per *active* user, not total users, so a "small role" isn't confused with "a role that uses little"). The table next to it shows

the same numbers as a list, flagging roles that had no usage in the period or accessed models outside their configured allowlist.

04 · OVERVIEW

Adoption

How many people actually use the tool, and how often — separate from usage volume (which lives in Costs and the Dashboard).



DAU, WAU, MAU, activation rate, a daily active-users series, and adoption broken down by role.

DAU / WAU / MAU

Users with real consumption (not just sign-in) in the last day, 7 days, and 30 days — fixed windows, independent of the period selected on this screen.

ACTIVATION RATE


The share of enabled users who actually used the tool in the selected period.

ADOPTION BY ROLE

How many users of each role are enabled versus how many were active — quickly shows where adoption is low.

Costs

How much AI usage cost during the period, broken down by role, model, and area.



Total cost, highest-cost role, verified-price percentage, cost per day, by role, by model, and by area.

Two ways to compute cost

This is the only screen whose behavior changes depending on the environment's configuration — worth understanding both modes:

MODE	WHEN IT APPEARS	WHAT IT MEANS
Verified pipeline	When the environment has its own cost-apportioning collection configured	Real cost in dollars, with a "verified price" flag separating what came from a confirmed source from what's still an estimate from that same source.
Estimate	Default — no cost collection configured	Cost computed from the real token count multiplied by an internal price table (approximate, subject to review). Still real data — only the dollar figure is estimated.

The screen always states, right below the title, which of the two modes is active — no guessing required.

i "Cost by area" only appears when the verified pipeline is configured, because area/cost-center isn't something LibreChat stores natively — it only exists if an external source provides it. See section 15 (Optional integrations).

Use Cases

What AI is being used for — today, measured indirectly, through the corporate Agent used.

Monitor LibreChat

OVERVIEW

Dashboard

Adaptation

Catche

Use Cases

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Online

Authorized Resources

PLATFORM

Audit Trail

Settings

Operational Status

<< Collapse

Monitor > Use Cases

EN PT admin AD

Use Cases Purpose and use-case category distribution

Showing example (synthetic) data — no connection to the LibreChat MongoDB.

The classification of **purpose** (professional / personal / undefined) and **use-case category** depends on the Classification Engine, planned for Phase 2. Classifying this requires analyzing conversation content with a model — a decision that requires prior approval from Security, Governance, and the DPO, and has not been implemented.

In the meantime, the available proxy for "use case" is the corporate Agent used, shown below — it indicates the domain of the query without reading the content.

USAGE BY CORPORATE AGENT (USE-CASE PROXY)

Agent	Uses	Share
Sales Insights	64	63%
Marketing Insights	37	37%

OVERALL VOLUME

Conversations	Prompts	Input tokens	Output tokens
231	842	512,000	298,000

Usage by corporate Agent (as a proxy for "use case") and overall volume of conversations, prompts, and tokens.

! Classifying each conversation's real purpose (professional, personal, by category) would require analyzing message content with an AI model — a decision that needs prior approval from Security, Governance, and whoever owns data privacy, which this version of the Monitor doesn't implement. The Agent used is the best available proxy today: it indicates the domain of the query without needing to read the content.

Maturity

A reference model for assessing how mature AI adoption is in the organization — today, documental: the criteria and levels are defined, but automatic score calculation isn't implemented yet.

GOVERNANCE

Monitor LibreChat

OVERVIEW

Dashboard

Adoption

Costs

Usage Cases

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Online

Authorized Resources

PLATFORM

Audit Trail

Settings

Operational Status

« Collapse

Monitored corporate AI usage

Monitor > Maturity

EN PT admin AD

Maturity AI Maturity Index — criteria and levels (Phase 2)

The maturity index calculation **is not implemented**. Three of the seven dimensions (prompt quality, reuse, contribution) require analysis of conversation content — which depends on the Classification Engine and prior approval from Security, Governance, and DPO. We prefer to keep the screen transparent about this rather than display a number that could not be explained or defended.

The four dimensions marked as available can already be measured with current data (usage metadata, Agents, MCP, and security detections) once the phase is approved.

PROPOSED DIMENSIONS AND WEIGHTS

Dimension	Weight	Evidence	Data available today
Usage regularity	15%	Consistent use, not just occasional.	yes
Prompt quality/structure	25%	Context, objective, format, constraints.	no — requires Phase 2
Use of corporate Agents	15%	Adoption of standardized resources.	yes
Use of authorized tools/MCP	10%	Use of tools with governance.	yes
Reuse/automation	15%	Reusable templates, flows, and assets.	no — requires Phase 2
Contribution/multiplication	10%	Sharing of practices and assets.	no — requires Phase 2
Good practices and security	10%	Use compliant with policies.	yes

MATURITY LEVELS (N0-N5)

Level	Name	Description
N0	Not used	No relevant evidence of use in the period.
N1	Explorer	Simple requests, little context, and low reuse.
N2	User	Frequent use with context and useful results.
N3	Advanced	Structured prompts, appropriate choice of models/Agents.
N4	Specialist	Integrates AI into processes and uses authorized tools/automation.
N5	Multiplier	Creates standards, Agents, templates, and disseminates good practices.

The seven proposed dimensions (with weight and expected evidence) and the six maturity levels, from "Not used" to "Multiplier".

The screen is transparent about the limitation: showing a calculated number without the criteria actually being measured would be worse than showing nothing.

Security & Risk

The Monitor's security detector, running in **shadow mode**: it detects and logs, but never blocks anything.

08 • GOVERNANCE

Monitor LibreChat

OVERVIEW

Dashboard

Adaptation

Costs

Usage Cases

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Entities

Authorized Resources

PLATFORM

Audit Trail

Settings

Operational Status

« Collapse

Monitored corporate AI usage

Monitor > Security & Risk

EN PT admin AD

Security & Risk Mode SHADOW — 842 messages analyzed in the period

Showing example (synthetic) data — no connection to the LibreChat MongoDB.

Operating in **SHADOW**: policies detect and log, but **nothing is blocked** — the Monitor is observational and does not intercept the chat. Only deterministic detections are active (credentials, private key, national ID, model outside the allowlist). Policies that depend on semantic classification appear as "awaiting Phase 2" on the Policies screen.

CRITICAL
1

HIGH
1

MEDIUM
2

LOW
0

INFORMATIONAL
0

DETECTIONS IN THE PERIOD (4)

Date	Severity	Policy	User	Evidence (masked)	Decision
9/20/2026, 2:18:09 PM	Critical	SEC-001 Credential or API key exposure	USR-EXAMPLE01	sk-F*****0012	ALLOW
9/20/2026, 2:18:09 PM	High	GOV-001 Model outside the role's allowlist	USR-EXAMPLE03	model "opus-5" outside the basic-user role allowlist	ALLOW
9/20/2026, 2:18:09 PM	Medium	PRIV-001 National ID or personal identifier	USR-EXAMPLE02	123.***-00	ALLOW
9/20/2026, 2:18:09 PM	Medium	PRIV-001 National ID or personal identifier	USR-EXAMPLE04	987.***-11	ALLOW

Detection counts by severity and the period's detection list, with evidence always masked.

What's detected today

Five policies actually run, all through local pattern matching — no content is ever sent to an AI model for this analysis, and none of the original text is stored:

- Credential or API-key exposure
- Explicit password or token in the conversation
- Private-key material
- Personal identifier (national-ID pattern)

- **Model usage outside a role's allowlist** (compared against LibreChat's real configuration)

The rest of the catalog (see section 10, Policies) is registered for traceability but doesn't produce detections yet — it depends on semantic content classification, a future phase.

! No policy is in blocking (enforcement) mode. Every row in the "Decision" column shows **ALLOW** — the Monitor is observing, not intercepting.

Alerts / Cases

The same detections from the Security screen, presented as a filterable alert board — built for case-by-case investigation.

Monitor LibreChat

OVERVIEW

Dashboard

Adaptation

Crate

Like Cases

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Online

Authorized Resources

PLATFORM

Audit Trail

Settings

Operational Status

<< Collapse

Monitored corporate AI usage

Monitor > Alerts / Cases

EN PT admin AD

Alerts / Cases Board of alerts generated by policies in shadow mode

Showing example (synthetic) data — no connection to the LibreChat MongoDB.

Alerts are **read-only** in this phase. The case management workflow (triage, assignment, comments, classification as false positive, closure) requires its own persistence and RBAC, planned for phase P2 — until then, handling of an alert happens outside the portal, through the current Security process.

ALERTS IN PERIOD
4

CRITICAL + HIGH
2
Triage priority

MESSAGES SCANNED
842

MODE
SHADOW
Nothing here sets isBlocked

FILTERS

Severity: All Policy: All

ALERTS (4)

Date	Severity	Policy	User	Evidence (masked)	Decision	Status
9/20/2026, 2:18:30 PM	Critical	SEC-001 Credential or API key exposure	USR-EXAMPLE01	sk-F*****0012	ALLOW	NEW
9/20/2026, 2:18:30 PM	High	GOV-001 Model outside the role's allowlist	USR-EXAMPLE03	model "opus-5" outside the basic-user role allowlist	ALLOW	NEW
9/20/2026, 2:18:30 PM	Medium	PRIV-001 National ID or personal identifier	USR-EXAMPLE02	123.***-00	ALLOW	NEW
9/20/2026, 2:18:30 PM	Medium	PRIV-001 National ID or personal identifier	USR-EXAMPLE04	987.***-11	ALLOW	NEW

Filters by severity and by policy, and the alert table with status.

This screen is read-only in this version — a full triage workflow (claiming a case, marking a false positive, closing it) would need its own access control, not yet implemented. For now, handling an alert happens outside the portal, through the organization's existing security process.

Policies

The complete catalog of planned AI governance policies, versioned as a file — not as a database record.

M

GOVERNANCE

Monitor LibreChat

OVERVIEW

Dashboard

Adoption

Costs

Health Factors

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Policiac

Authorized Decisions

PLATFORM

Audit Trail

Settings

Operational Status

Monitor > Policiac

EN PT admin AD

Policiacs

Versioned catalog at policies/*.yaml in the repository

The full blueprint catalog is registered, but only **deterministic** detection policies are actually running. The rest depend on the Classification Engine (Phase 2) or a synchronous Guardian (Phase 3), and are listed here for traceability — they don't produce detections today.

POICIES IN THE CATALOG

20

ACTIVE (DETERMINISTIC)

5

AWAITING FUTURE PHASE

15

ACTIVE — RUNNING IN SHADOW (5)

ID	Name	Domain	Severity	Mode	Detection	Note
GOV-001	Model outside the role's allowlist	governance	High	SHADOW	allowlist	Compares the model used in transactions against the per-role model allowlist (Authorized Resources section). The blueprint suggests enforcement, but the Monitor is observational: it detects after the fact and does not block usage. Enforcement would require a synchronous inline gate (a later phase).
PRIV-001	National ID or personal identifier	privacy	Medium	SHADOW	dip	Detects a national-ID numeric pattern (000.000.000-00 or 000000000000).
SEC-001	Credential or API security key exposure		Critical	SHADOW	regex/secret-detector	Detects API key patterns (OpenAI sk-, AWS AKIA, GitHub ghp-, Google Alza) in message text.
SEC-002	Explicit password security or token		Critical	SHADOW	regex/contexto	Detects keywords (password, token, bearer) near a value in the same passage.
SEC-003	Private key or certificate	security	Critical	SHADOW	pattern-detector	Detects private-key/certificate headers (-----BEGIN ... PRIVATE KEY-----).

REGISTERED, AWAITING FUTURE PHASE (15)

ID	Name	Domain	Severity	Mode	Detection	Note
BEH-001	Content unsuitable for a corporate environment	behavior	Medium	SHADOW	classifier	Requires a semantic behaviour classifier — planned for a later phase.
BEH-002	Harassment or direct threat	behavior	High	SHADOW	classifier	Requires a semantic classifier — planned for a later phase. Must never trigger an automatic disciplinary decision (ADR-008).
COST-001	Consumption above the expected baseline	cost	Medium	ALERT	threshold/trend	Per-role cost threshold not yet calibrated — see the Costs dashboard for the raw figures.
DATA-001	Classified confidential content	data	High	SHADOW	marker+classifier	Requires a semantic classifier and document sensitivity tagging — planned for a later phase.
GOV-002	Agent outside the role's allowlist	governance	Medium	SHADOW	allowlist	Requires an agent-to-role ACL mapping not yet integrated — planned for a later phase.
GOV-003	MCP integration outside the role's allowlist	governance	High	SHADOW	allowlist	Requires correlating MCP log entries with the user's role — planned for a later phase.
GOV-004	Public sharing not permitted	governance	High	SHADOW	permission/event	Depends on a LibreChat sharing event not captured in this build — planned for a later phase.
OPS-001	Inline gate latency above SLO	ops	Medium	ALERT	metric	There is no synchronous inline gate in this build (read-only) — not applicable until that phase ships.
OPS-002	Event publishing failure	ops	High	ALERT	metric/dlq	There is no event pipeline in this build (ADR-002) — not applicable.
OPS-003	Classifier unavailable	ops	Medium	FAIL-OPEN	health	There is no classification engine in this build — not applicable.
PRIV-002	Personal data in bulk	privacy	High	SHADOW	dip+volume	Requires detecting volume/batches of personal records — awaits a classification engine (a later phase).
PRIV-003	Sensitive personal data	privacy	High	SHADOW	dip+classifier	Requires a semantic classifier for sensitive categories (health, religion, etc.) — planned for a later phase.
USE-001	Professional use	usage	Info	LOG	classifier	Classifying professional vs. personal intent requires a classification engine — planned for a later phase.
USE-002	Occasional personal use	usage	Low	LOG	classifier	Same as USE-001 — planned for a later phase.
USE-003	Recurring personal use	usage	Low	SHADOW	trend	Depends on USE-002 being classified — planned for a later phase.

« Collapse

Monitored corporate AI usage

Catalog of 20 policies: the ones actually running (shadow mode) and the ones awaiting a future phase.

- ✓ Each policy lives in its own `policy.yaml` file inside the policy directory. Adding or adjusting a policy means editing that file — no code change or rebuild required.

Authorized Resources

Which models, Agents, and integrations each role may use — the reference the "model outside the allowlist" policy compares against.

Monitor LibreChat

OVERVIEW

Dashboard

Admission

Costs

Usage Cases

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Online

Authorized Resources

PLATFORM

Audit Trail

Settings

Operational Status

« Collapse

Monitored corporate AI usage

Monitor > Authorized Resources

EN PT admin AD

Authorized Resources Allowlist of models, Agents and MCPs by profile

The allowlist is read **live** from db.configs in LibreChat (overrides.modelSpecs), which is the same source used by chat — there is no manual mirror to go stale. The GOV-001 policy compares actual usage against this list, and only raises an alert for profiles with **enforce** turned on.

User count unavailable ("—"): no connection to the LibreChat Mongo. The allowlist below is still valid — it comes from configuration, not from the database.

MODELS AND MCP BY PROFILE

Profile	Users	Allowed models	Enforce	MCP
CORPORATE AGENTS (0)				
No Agents returned — no connection to the LibreChat Mongo.				
		MCP INTEGRATIONS		
		MCP	Web search	
		Example MCP Integration	Disabled	

Model allowlist by role (read live from LibreChat's own configuration), corporate Agents, and MCP integrations.

The model allowlist isn't the Monitor's own record — it's read directly from LibreChat's configuration, the same source the chat itself uses. That keeps the portal from holding a stale copy the moment someone changes the configuration through LibreChat.

Audit Trail

Two distinct audit trails, on the same screen: who accessed the Monitor itself, and who used the corporate integrations via MCP.

« Collapse

Monitored corporate AI usage

9/20/2026, 2:15:58 PM	GET	/api/v1/resources	200	0 ms
9/20/2026, 2:15:58 PM	GET	/api/v1/analytics/cost	200	0 ms
9/20/2026, 2:15:58 PM	GET	/api/v1/analytics/usage	200	1 ms
9/20/2026, 2:15:58 PM	GET	/api/v1/analytics/adoption	200	1 ms
9/20/2026, 2:14:58 PM	GET	/api/v1/analytics/conduct	200	1 ms
9/20/2026, 2:14:58 PM	GET	/api/v1/analytics/security	200	2 ms


Corporate MCP queries and the Monitor's own access history — method, route, status, and duration for each call.

- !
- The Monitor's own access history lives in memory — it's lost when the service restarts.

Long-term persistence arrives together with more granular access control, not yet implemented in this version.

Settings

A reference for what's configured today, and for how role-based access control will work once it's ready — for now, documentaryal.



Planned access roles (RBAC) and where each system parameter comes from.

In this version, the Monitor has a single shared login — no differentiated access roles. The RBAC design shown here is the target for once dedicated authentication is implemented; don't rely on this screen as effective access control today.

Operational Status

The Monitor's own health: is it connected to MongoDB? does the data match the expected shape? which optional integrations are turned on?

Monitor LibreChat

OVERVIEW

Dashboard

Adoption

Costs

Usage Cases

Maturity

GOVERNANCE

Security & Risk

Alerts / Cases

Online

Authorized Resources

PLATFORM

Audit Trail

Settings

Operational Status

<< Collapse

Monitored corporate AI usage

Monitor > Operational Status

EN PT admin AD ↗

Operational Status Health of Monitor components and SLOs

OVERALL STATUS
OK

LIBRECHAT MONGODB
not configured
Endpoints return synthetic data

VERSION
0.1.0

UPTIME
6min

SLOS

SLO	Target	Note
Monitor API availability	>= 99.9%	To be calibrated after the first month in production.
Query latency p95	<= 1500 ms	Aggregate queries against the LibreChat database.
Inline pre-check latency p95	<= 150 ms	Not applicable: this build is observational only and has no inline enforcement (ADR-001).
correlation_id coverage	>= 99.5%	Not applicable: no event pipeline in this build (ADR-002).

OPTIONAL INTEGRATIONS

Every one of these is off by default and only turns on when a deployment specifically configures it — the core dashboards (adoption, usage, security, audit) don't need any of them.

Feature	Status	Details
Cost pipeline	Off	COST_COLLECTION_NAME is not set — Costs uses a token-times-price estimate.
MCP audit	Off	GOOGLE_CLOUD_PROJECT and/or MCP_AUDIT_SERVICE_NAMES aren't set — the MCP section of Audit Trail uses synthetic data.
MCP integrations	Enabled	Still using the placeholder example in policies/mcp-integrations.yaml — replace it with your own integrations.

Overall status, MongoDB connection, availability targets (SLOs), and the optional integrations panel.

This is the most useful screen for whoever is setting up the Monitor for the first time — it answers, on its own, "what will work with my environment" without needing to read any config file. See the next section for detail on each optional integration.

Optional integrations

The Monitor was built to work with real data against any LibreChat instance, with no configuration beyond the database connection. Three features, however, are optional — they ship off by default and only turn on once someone configures them explicitly:

FEATURE	OFF (DEFAULT)	ON
Cost pipeline	Costs uses a token × price estimate	Costs uses an external, already-apportioned cost collection
MCP audit	The MCP section of Audit Trail shows example data	Reads the real history of configured MCP integrations
MCP integrations catalog	No integration registered in Authorized Resources	Shows the integrations and which roles may use them

None of the three requires touching code — they're config files and environment variables. Whoever installs the Monitor should consult the repository's technical documentation (README) for the exact steps for each.

FAQ

Are the numbers I'm seeing real?

If the yellow "Showing example (synthetic) data" banner appears, no — it's a demonstration, with no connection to LibreChat's MongoDB. Without that banner, the numbers are real.

Why did cost change differently between two visits?

Check whether the Costs screen is in "estimate" or "verified pipeline" mode (section 05) — the two use different sources and aren't directly comparable period over period if the configuration changed in between.

Can I use the Monitor to block a user or a prompt?

No. The Monitor is observational on every screen — it never intercepts or prevents a conversation. Any blocking action has to happen in LibreChat's own configuration.

I forgot the Monitor's login password.

Talk to whoever administers the environment — password reset happens outside the portal, directly in the service's configuration.